

LeWAM: END-TO-END WORLD AND ACTION MODELING WITH JEPAS

Minghao Fu^{*1} Tavis Siebert^{*2} Eryk Halicki³ Randall Balestrierio^{4,5}
¹UC San Diego ²ETH Zurich ³University of British Columbia ⁴Brown University ⁵AMI Labs



Website



Code

ABSTRACT

Joint-embedding predictive architectures (JEPAs) learn compact latent world models without pixel reconstruction. However, existing JEPAs typically learn representations and dynamics independently of downstream control objectives, complicating the pipeline and leaving unexplored how world- and action- learning benefit each other. We introduce **LeWAM**, an end-to-end JEPA for control with: (1) a joint training objective over next latent state prediction and action flow matching, and (2) an isotropic Gaussian constraint on the latent state, enabling a single model to support direct policy execution and efficient planning by evaluating sampled actions with its learned latent dynamics. Compared with prior JEPAs in visual planning, LeWAM improves success rate from 26.8% to 89.7% on contact-rich tasks and from 10.9% to 46.2% on long-horizon tasks, and plans $32.3\times$ faster. On behavior cloning tasks in simulation and physical robots, LeWAM is also competitive with established methods. Interpretability analyses show that LeWAM learns control-relevant representations and action distributions effectively guiding test-time planning and control.




Figure 1: LeWAM Training. A shared encoder encodes observed frames with SIGReg (Balestrierio & LeCun, 2025) regularization. The Mixture-of-Transformers (Liang et al., 2025) learns latent-state prediction from clean actions and action flow matching from noisy actions with goal information. **Results Overview.** The mean of (a) planning success on contact-rich tasks; (b) planning success on a 100-step goal offset; (c) plans/s versus CEM (Rubinstein, 1999); (d) action chunks per TFLOP versus $\pi_{0.5}$ (Physical Intelligence et al., 2025). Full results are in Section 4.

^{*}Equal contribution. Correspondence to m9fu@ucsd.edu.

LeWM, V-JEPA 2	SIGReg, DINOv2	DreamZero, Fast-WAM	CEM, MPPI, iCEM	Navigation, 2D Action
✗ Separated stage for control	✗ Strong semantic representation	✗ Generative Video WAMs	✗ Search from broad priors	✗ Short-horizon simple tasks
LeWAM: End-to-End World and Action Modeling with JEPAs				
✓ Unified world model and control	✓ Control-relevant representation	✓ No reconstruction ✓ No pre-training	✓ Policy-guided search ✓ Gradient-based refinement	✓ Contact-rich ✓ Long horizon

Figure 2: **Positioning of LeWAM.** LeWAM addresses key limitations in prior visual control frameworks: (1) Unlike JEPAs requiring a separate stage for control, it unifies world modeling and action policy learning end-to-end; (2) Compared to learning strong semantic representations, action prediction tailors representations to be more control-relevant; (3) In contrast to pixel-generative world action models, it operates in latent space without reconstruction; and (4) Rather than searching from broad priors, it guides planning using learned policy proposals with latent gradient refinement.

1 INTRODUCTION

Visual control requires autonomous agents to complete tasks by selecting actions based on visual observations of the world. World models address this challenge by predicting the consequences of candidate actions, allowing agents to anticipate how the world will evolve before executing them. Joint-embedding predictive architectures (JEPAs) (LeCun, 2022) offer a compact world modeling paradigm where future world states are predicted in latent space rather than by reconstructing pixels. Recent works such as DINO-WM (Zhou et al., 2024) and LeWorldModel (LeWM) (Maes et al., 2026b) demonstrate the promise of applying action-conditioned latent world models to downstream control. In world model pretraining, these pipelines use pretrained visual embeddings (Zhou et al., 2024; Assran et al., 2025) or distributional regularizers such as SIGReg (Balestriero & LeCun, 2025) to learn strong semantic representations and prevent collapse. However, these methods alone do not reliably retain control-relevant information; JEPAs can favor temporally persistent distractors (Sobal et al., 2022), while action-conditioned dynamics may retain extraneous observation information (Brandfonbrener et al., 2023).

Control with these models typically requires an online solver such as the cross-entropy method (CEM) (Rubinstein, 1999) or a policy model post-trained in a separate stage, such as GC-IDM (Nguyen et al., 2026) and PRISM (Wang et al., 2026). When initialized with a broad Gaussian prior over action sequences, CEM demands extensive rollouts to identify viable candidates, particularly as action dimensionality or planning horizon grows (Huang et al., 2021). Contact transitions also complicate local action optimization because the dynamics can switch between modes (Pang et al., 2023). Post-training a action prediction as planner over frozen JEPA embeddings can resolve them, but leaves action supervision disconnected from visual representation learning, and specialized to the action demonstrations.

We introduce **LeWAM**, an end-to-end JEPA that jointly learns representation, action-conditioned dynamics, and a stochastic policy in a single training stage (Figure 1). Training combines only two ingredients: a joint training objective over next-latent-state prediction and action flow matching, and SIGReg (Balestriero & LeCun, 2025) to regularize the latent states. A Mixture-of-Transformers (MoT) (Liang et al., 2025) predictor neatly satisfies action-conditioned dynamics modeling over clean actions and policy learning over noisy actions through flow matching, jointly shaping representations to preserve directions relevant for predicting actions and anticipating their effects. Consequently, LeWAM supports both direct policy execution and visual planning by refining action sequences sampled from its learned policy. Our contributions are threefold:

- The first JEPA that jointly learns representation, action-conditioned latent dynamics, and stochastic policy, using one simple joint prediction objective, one hyperparameter for latent space regularization, and just 17M parameters.
- A planning procedure that jointly imagines the consequences of actions and how to act in the resulting states, which samples candidate action sequences from the learned action distribution and refines them with gradient descent through learned dynamics.

- Across ten heterogeneous simulated environments, LeWAM achieves faster planning speedup over LeWM (Maes et al., 2026b) alongside consistent success rate improvements, particularly on contact-rich and long-horizon tasks. On goal-agnostic task completion in simulation and physical robot execution, LeWAM performs competitively with established baselines. Mechanistically, interpretability analyses show that LeWAM learns control-relevant representations and action-conditioned dynamics, while its structured action distributions yield feasible proposals that guide test-time planning.

2 RELATED WORK

World Models. World models (Ha & Schmidhuber, 2018) learn environment dynamics to enable agents to simulate candidate action outcomes for planning and control. Early approaches like PlaNet (Hafner et al., 2019) perform planning via latent rollouts, whereas TD-MPC (Hansen et al., 2022) and TD-MPC2 (Hansen et al., 2024) ground latent dynamics in reward and value predictions. Dreamer (Hafner et al., 2020; 2023) learns generative latent dynamics with pixel reconstruction, while VideoGPT (Yan et al., 2021) and IRIS (Micheli et al., 2023) model discrete visual tokens. Within this landscape, LeWAM explicitly targets latent dynamics and policy representations directly from offline trajectories, without reward supervision or pixel reconstruction.

Joint-Embedding Predictive Architectures. JEPAs predict target states in representation space without pixel reconstruction, using variance and covariance constraints (Bardes et al., 2022), momentum targets (Assran et al., 2023; Bardes et al., 2024), or distributional regularization (Balestriero & LeCun, 2025; Wu et al., 2026) to prevent collapse. For action-conditioned world modeling, DINO-WM (Zhou et al., 2024) and V-JEPA 2-AC (Assran et al., 2025) build on pretrained representations, whereas LeWM (Maes et al., 2026b) leverages SIGReg (Balestriero & LeCun, 2025) to learn the encoder and latent dynamics from scratch. Additional action prediction objectives attempt to make representations action-aware in PLDM (Sobal et al., 2025) or as an anti-collapse signal in SMWM (Ivashkov et al., 2026). LeWAM jointly trains action-conditioned latent dynamics and policy in one predictor, allowing representations to preserve more control-relevant directions.

Control with Latent Dynamics. Conventional planning in latent world models uses population-based solvers, such as CEM (Rubinstein, 1999), iCEM (Pinneri et al., 2021), MPPI (Williams et al., 2017), or gradient-based methods (Henaff et al., 2017; Huang et al., 2021). Population-based search can require many model rollouts as the action dimension or horizon grows (Pinneri et al., 2021; Huang et al., 2021), and contact mode changes in manipulations also complicate local optimization (Pang et al., 2023) (more discussion in Appendix E). GC-IDM (Nguyen et al., 2026) post-trains a goal-conditioned action predictor for direct control, whereas PRISM (Wang et al., 2026) trains a Gaussian prior head on frozen world-model features – its predicted mean and standard deviation guide MPPI sampling. LeWAM jointly trains the latent world model and a flow-based action predictor, using its sampled actions for direct control or gradient-based planning.

World Action Models. These models couple world prediction with action prediction (Table 12, Appendix E). Built on pretrained video generative models, DreamZero (Ye et al., 2026) jointly models future video and actions, while Fast-WAM (Yuan et al., 2026) retains video co-training but skips future video generation at inference. UVA (Li et al., 2025), UWM (Zhu et al., 2025), and Cosmos 3 (NVIDIA, 2026) specifically support both forward dynamics modeling (FDM) and policy learning, but optimize only one objective per sample, rather than both jointly on the same pass. Furthermore, explorations on how their methods benefit planning remain limited. Instead of borrowing costly video generation, LeWAM jointly trains its encoder and predictor end-to-end without pixel reconstruction. Each training sample supplies clean actions for FDM and noisy actions for stochastic policy modeling. Thus, LeWAM can learn dynamics from both expert and non-expert interactions, and specifically allows action-conditioned dynamics to evaluate and refine actions at test time.

3 METHODOLOGY

Given reward-free offline trajectories of pixel observations $\mathbf{o}_{1:T}$ and actions $\mathbf{a}_{1:T}$, we jointly train a visual encoder and a predictor including action-conditioned dynamics and a flow matching policy, supporting direct control and planning with action-conditioned latent dynamics.

3.1 END-TO-END TRAINING WITH LEWAM

Architecture. Figure 1 summarizes the model at a high level. Let $\mathbf{z}_t = \text{Enc}(\mathbf{o}_t)$ denote the latent representation of observations $\mathbf{o}_t$. All time indices refer to raw environment steps, such that action $\mathbf{a}_t$ transitions $\mathbf{o}_t$ to $\mathbf{o}_{t+1}$. We use a frameskip of K environment steps, so future observations used as dynamics targets are spaced at K -step intervals. In contrast, the policy models a contiguous action chunk $\mathbf{A}_{t,C} = (\mathbf{a}_t, \dots, \mathbf{a}_{t+C-1})$ with $C \geq K$. For visual planning, we assume access to goal information $\mathbf{c}_t$ that contains the goal image embedding $\mathbf{z}_g$ and its normalized remaining horizon from $\mathbf{z}_t$.

The encoder consists of a ResNet-18 backbone (He et al., 2016), spatial softmax pooling, and a multilayer perceptron (MLP) with two projection layers. The same encoder processes all the pixel observations: history, future target, and goal image. The predictor is comprised of a state branch f_z and an action branch f_a . Specifically,

$$\begin{aligned}\hat{\mathbf{z}}_{t+K} &= f_z(\mathbf{q}_{t+K}^z; \mathbf{z}_{\leq t}, \mathbf{A}_{t,K}), \\ \hat{\mathbf{v}}_{t,C} &= f_a(\mathbf{A}_{t,C}^{(s)}; \mathbf{s}, \mathbf{z}_{\leq t}, \mathbf{c}_t)\end{aligned}\quad (1)$$

where $\mathbf{A}_{t,C}^{(s)}$ is a noisy C -action chunk at flow timestep $s \in [0, 1]$ (see Equation 2) and $\mathbf{q}_{t+K}^z$ is a learned query token specifying the future state to predict. Integrating the predicted velocity field $\hat{\mathbf{v}}_{t,C}$ yields the predicted action chunk $\hat{\mathbf{A}}_{t,C}$.

To simultaneously balance latent-prediction and flow-matching objectives while allowing f_z and f_a to share information about states and actions, we use the Mixture-of-Transformers (MoT) (Liang et al., 2025) architecture. That is, each branch applies independent query, key, value, and feed-forward projections for its respective modality (latent states or actions) while sharing one masked self-attention. Figure 3 specifies the information available to each branch during attention. The state query $\mathbf{q}_{t+K}^z$ reads history latents and clean actions, and noisy action tokens $\mathbf{A}_{t,C}^{(s)}$ read history latents and their own causal prefix. History latents attend only to themselves, blocking indirect clean-action leakage to f_a . Goal information conditions only the noisy action tokens, leaving the dynamics unconditioned on either. Appendix A gives the full token-level mask, layer sizes, and training settings.

Training Objective. For flow matching (Lipman et al., 2023), given a recorded action chunk $\mathbf{A}_{t,C}$, sample s uniformly from $[0, 1]$ and $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$, then construct

$$\mathbf{A}_{t,C}^{(s)} = (1-s)\epsilon + s\mathbf{A}_{t,C}, \quad \mathbf{v}_{t,C} = \frac{d\mathbf{A}_{t,C}^{(s)}}{ds} = \mathbf{A}_{t,C} - \epsilon \quad (2)$$

The action branch regresses the path velocity $\mathbf{A}_{t,C} - \epsilon$ from $\mathbf{A}_{t,C}^{(s)}$. The latent dynamics and action flow matching losses use mean squared error (MSE), averaged over the scalar coordinates of each prediction:

$$\begin{aligned}\mathcal{L}_{\text{dyn}} &= \mathbb{E}_{\mathcal{D}}[\text{MSE}(\hat{\mathbf{z}}_{t+K}, \mathbf{z}_{t+K})], \\ \mathcal{L}_{\text{act}} &= \mathbb{E}_{\mathcal{D}, s, \epsilon}[\text{MSE}(\hat{\mathbf{v}}_{t,C}, \mathbf{A}_{t,C} - \epsilon)].\end{aligned}\quad (3)$$

We use SIGReg (Balestriero & LeCun, 2025) to preserve latent diversity through an isotropic Gaussian distributional constraint. Let $\mathbf{Z}$ collect the encoded history and target observations in a mini-batch. With M random unit directions $\mathbf{u}^{(m)}$ and the univariate Epps-Pulley (Epps & Pulley, 1983) test statistic T against $\mathcal{N}(0, 1)$, the regularizer is

$$\mathcal{L}_{\text{sig}} = \frac{1}{M} \sum_{m=1}^M T(\mathbf{Z}\mathbf{u}^{(m)}). \quad (4)$$

With regularization weight λ_{sig} , the full objective is

$$\mathcal{L}_{\text{LeWAM}} = \mathcal{L}_{\text{dyn}} + \mathcal{L}_{\text{act}} + \lambda_{\text{sig}}\mathcal{L}_{\text{sig}}. \quad (5)$$




Figure 3: **MoT Attention.** Filled cells permit attention, while white cells are masked. The state query reads history and clean actions, and noisy action tokens read state history and their own causal prefix. Small grids show causal attention. Arrows mark the state branch and action velocity branch.




Figure 4: **Planning with LeWAM.** Each MoT predictor stage samples N action chunks from Gaussian noise and uses the first K actions in each chunk, denoted $A_{t,K}$, to advance the latent dynamics from step t to $t + K$. Five-point curves depict action chunks (we use $K = 5$), and stacked curves denote N sampled plans of length L . Step 1 refines all plans by gradient descent through the frozen LeWAM dynamics using the goal-state cost. Step 2 selects the lowest-cost refined iterate as the final plan, shown as A^* . Appendix A.1 gives the full procedure in a pseudo-algorithm.

3.2 EFFICIENT CONTROL WITH LEWAM

LeWAM supports both direct control through its learned policy and planning through latent state-action joint imaginations.

Direct Control. With history and optional goal context fixed, one C -action chunk is sampled by integrating along the learned velocity field over the interval $[0, 1]$ from Gaussian noise:

$$\frac{dA_{t,C}^{(s)}}{ds} = f_a(A_{t,C}^{(s)}; s, z_{\leq t}, c_t), \quad A_{t,C}^{(0)} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}), \quad \hat{A}_{t,C} = A_{t,C}^{(1)} \quad (6)$$

The first K actions are executed and resampling begins from the resulting observation.

Planning. LeWAM uses a gradient-based refinement strategy over a population of action candidates sampled from the learned action distribution (Figure 4). To simplify notation, let $A = (a_t, \dots, a_{t+L-1})$ denote a plan to reach the goal z_g after L raw environment steps from time t , where L is a multiple of frameskip K .

Starting from the observed history $z_{\leq t}$, one population of N candidate plans is generated from the learned action distribution by alternating action sampling and latent dynamics. Specifically, at each step, Eq. (6) generates N action chunks from the observation history and goal. The forward dynamics branch uses the first K actions of each chunk to predict the N corresponding next states. Each predicted state is appended to the fixed-length history of its candidate, the planning horizon is updated, and the process continues for L total steps. This leaves N action sequence candidates $\{A^{(n)}\}_{n=1}^N$. Temporally coherent policy rollouts provide structured starting points from which we assume that nearby improvements can be found via gradient descent through the latent dynamics (Henaff et al., 2017; Jyothir S V et al., 2023). Starting from $\{A^{(n)}\}_{n=1}^N$, we optimize each candidate with respect to the latent cost-to-goal

$$\mathcal{C}(A) = \text{MSE}(\hat{z}_{t+L}(A), z_g) \quad (7)$$

With learning rate η , the goal-cost update takes the gradient-descent form

$$a_{t+k} \leftarrow a_{t+k} - \eta \frac{\partial \mathcal{C}}{\partial a_{t+k}}, \quad k = 0, \dots, L-1, \quad (8)$$

The descent iterate with the lowest $\mathcal{C}$ across all candidates is chosen. In practice, we add a small proximity penalty to discourage large changes from the sampled plans (Appendix A.1). To balance efficiency and accuracy, the first K actions of a plan are executed, and the entire two-stage process repeats until the goal is reached, as in receding-horizon model predictive control.

4 EXPERIMENTS

Our experiments aim to address the following questions: (1) Can a single LeWAM architecture learn stably from offline data and support visual planning across heterogeneous tasks and settings?




Figure 5: **Simulation Environments.** Symbols beneath each frame denote overlapping task attributes, with their meanings shown in the legend to the left.

(2) Does LeWAM jointly learn informative action-conditioned latent dynamics, control-relevant representations, and multi-modal distributions over actions? (3) Can LeWAM further support behavior cloning in simulated environments and physical-robot control?

4.1 EVALUATION PROTOCOL

Heterogeneous Simulated Environments We evaluate control across ten simulated navigation and manipulation environments (Figure 5). To ensure the evaluation is reliable, the environments vary in task functionality, temporal horizon, and collection regime, which cover expert demonstrations, random-policy interactions, and mixed-quality data. Four environments adapt the benchmarks used in LeWM (Maes et al., 2026b). The other six add challenging contact-rich manipulation and navigation. Appendix B describes the tasks, data properties, construction procedure, and success criterion.

Evaluation Protocols. We distinguish two separate evaluation settings. Visual planning assumes access to a goal image, and a goal-conditioned policy or planner must reach the depicted state from its current (random) starting position within a fixed step budget. Task completion provides no goal image and evaluates behavior cloning from demonstrations: policies execute actions from start to finish to complete one specific task. Within each setting, methods share training and evaluation data, observations, action limits, goal offsets, execution budgets, and success predicates. Appendix B.2 details goal construction, evaluation conditions, and success criteria.

4.2 VISUAL PLANNING WITH LEWAM

We compare with four goal-conditioned baselines spanning latent-dynamics planning, an inverse-dynamics policy on a frozen world model, and direct policy execution:

- **LeWM (CEM)** (Maes et al., 2026b). Learning compact action-conditioned latent dynamics end-to-end with SIGReg, planning with CEM over learned latent dynamics.
- **LeWM (IDM)** (Nguyen et al., 2026). Post-trains a goal-conditioned IDM on frozen LeWM.
- **DINO-WM** (Zhou et al., 2024). Learns dynamics on a frozen DINOv2 encoder with CEM planner.
- **GCBC** (Ghosh et al., 2021). Trains a goal-conditioned policy on frozen DINOv2 features.

Figure 6 reports goal-reaching success across ten environments. Across the ten tasks, LeWAM averages 93.5% success versus 46.9% for LeWM (CEM) (Maes et al., 2026b), a 46.6 percentage-point gap. LeWAM exceeds that baseline on every task, with particularly large gaps on contact-rich ToolHang (94.7 versus 35.4) and Transport (83.3 versus 2.8). On the logically structured Puzzle task, LeWAM reaches 80.7% success versus 44.0% for goal-conditioned GCBC (Ghosh et al., 2021), highlighting the advantages of dynamics-based planning over model-free execution.

Long-Horizon Planning. We evaluate planning robustness with the same checkpoints, over varying horizons using four manipulation tasks with diverse visual goals: Push-T, Cube, Scene, and Puzzle. Unlike simple navigation benchmarks, these tasks introduce significantly higher behavioral diversity as the planning horizon grows. As shown in Figure 7, LeWAM exhibits less performance degradation than LeWM with increasing horizon length, highlighting its stability in long-horizon planning.




Figure 6: **Visual Planning across Ten Simulated Environments.** Each panel reports the success rate under a goal offset of $L = 25$ environment steps and permits $2L = 50$ actions across 3 seeds. Details are present in Appendix B.2. Error bars denote one standard deviation; ★ indicates models with end-to-end training. Linear axes start at 50% for the first five panels and 0% elsewhere.




Figure 7: **Visual Planning Success across Different Horizons.** The task panels show mean success rates across 3 seeds. Each sweep starts at the first trajectory frame but with different goal offsets.

Planning Speed. Appendix Table 6 compares the time for one planning call on four visual-goal tasks. At the reported solver settings, LeWAM averages 52.50 plans/s across tasks, versus 1.62 plans/s for LeWM (CEM) (Maes et al., 2026b) with a $32.3\times$ acceleration, as shown in Figure 1(c).

4.3 UNDERSTANDING LEWAM’S DYNAMICS, REPRESENTATIONS, AND ACTIONS

Latent Representation Probing. To better measure what information is learned by our joint-training objective, we evaluate how well linear probes recover physical states, proprioceptive information, and actions from LeWAM’s representations. As highlighted in Table 1 and Appendix C.3, training with an action modeling objective noticeably improves LeWAM’s linear action and proprioception decodability compared to LeWM, increasing average proprioception, action, and IDM R^2 by +0.21, +0.20, +0.19 respectively (see Table 9, Appendix C.3).

The Role of SIGReg. Furthermore, while SIGReg (Balestriero & LeCun, 2025; Maes et al., 2026b) is primarily used in JEPAs to prevent representation collapse during training, we observe that in our setup, LeWAM’s joint prediction objective over latent states and actions is *sufficient* in preventing collapse without SIGReg or any other tricks such as variational constraints or stop-gradients. Table 7 highlights this finding; removing SIGReg makes state information unrecoverable in LeWM, while it is unnecessary for all LeWAM probes (see Appendix C.2 for more). *Nevertheless*, we observe that SIGReg itself is invaluable to planning. On Push-T, the planning gain is 14.7 percentage points with SIGReg, compared to 2.7 without it (Table 7). The probing results of Table 1 and Table 9 demonstrate that, beyond ensuring anti-collapse, SIGReg improves accessibility of object information in the latent state. We hypothesize that this translates to improved planning by making plan-selection metrics such as latent distance-to-goal more sensitive to changes in the underlying environment state.

Table 1: Linear probing and the role of SIGReg. All linear probes report R^2 score using ridge regression. *Object* denotes the average R^2 over all object ground truth information and *Proprio* the ground truth proprioceptive state. *Action* predicts the action chunk from latent z_t ; parenthesized IDM scores use both z_t, z_{t+K} where reported. *Eff. rank* denotes effective rank (Roy & Vetterli, 2007).

Environment	Model	SIGReg	Eff. rank	Object $\uparrow$	Proprio $\uparrow$	Action (IDM) $\uparrow$
Push-T	LeWAM	Yes	42.90	0.90	0.86	0.40(0.75)
		No	10.94	0.88	0.85	0.38(0.73)
	LeWAM (MSE)	No	6.96	0.66	0.86	0.36 (0.74)
		Yes	94.20	0.95	0.80	0.29(0.71)
	LeWM	No	16.23	0.00	0.00	0.00 (0.00)
		Yes				

Figure 8(a) compares full LeWAM with its no-SIGReg ablation. Removing SIGReg lowers planning success in the three tasks shown, most strongly on Push-T (25.3 percentage points), followed by Puzzle (5.3) and Scene (4.0).

Why Train the Visual Encoder End-to-End? Evaluated under the same downstream setup, joint encoder training improves direct control over a frozen DINOv2 (Oquab et al., 2023) encoder (Figure 8(b)). This suggests strong semantic representations are suboptimal for planning and control, and that action prediction can make representations policy-friendly. Furthermore, planning improvements indicate that in the presence of policy joint training, training the encoder from scratch still preserves useful action-conditioned dynamics. Appendix C.4 provides aligned visual reconstructions and recorded-action open-loop rollouts for four simulated tasks.

Why Model Stochastic Action Distributions? We contrast action flow matching with sampling from a Gaussian centered at a deterministic action prediction obtained by training our action branch with standard MSE loss. Under the same dynamics objective, flow matching yields larger planning gains across tasks (Figure 8(c)). In this comparison, the learned stochastic action model provides more effective planning proposals than Gaussian perturbations around a point prediction. Furthermore, flow matching retains more object information than a deterministic head ($R^2 = 0.69$ versus 0.20) and slightly more proprioceptive information (0.86 versus 0.84). While MSE targets the conditional mean $\mathbb{E}[A_{t,C} \mid z_{\leq t}, c_t]$, flow matching models the conditional action distribution, which may place greater demands on the latent representation to preserve information that distinguishes between different modes of behavior such as object-level differences. Appendix C.5 further demonstrates the ability of flow matching to learn multimodal action distributions in a two-door navigation problem.




Figure 8: (a) Planning gain (percentage points) from SIGReg. (b) Direct success and planning gain for visual encoders, averaged over four tasks. (c) Planning gain for flow (circles) and MSE (squares) action heads. (d) Task-completion success for LeWAM and Diffusion Policy.

4.4 BEHAVIOR CLONING AND PHYSICAL ROBOT CONTROL

While visual planning relies on goal images to optimize actions online, goal-agnostic tasks require executing behaviors directly. In this section, the robot must implicitly infer task progress and react to changing object poses and contacts without explicit goal conditioning. We train LeWAM without goal context c_t while retaining other objectives. At test time, the policy directly samples action chunks from observation history for execution, bypassing the test-time planning.




Figure 9: Paired scene and wrist view examples from three physical-robot tasks.

Task Completion in Simulation. We benchmark LeWAM on three simulated tasks requiring complex manipulation without goal prompts: DrawerCleanup, Transport, and ToolHang. Figure 8(d) compares task-completion success with Diffusion Policy (Chi et al., 2023). LeWAM achieves competitive performance on two tasks and improves success by 16 percentage points on Transport. These results show that LeWAM supports goal-agnostic behavior cloning alongside goal-conditioned planning. Per-task standard deviations appear in Appendix Table 10.

Physical-Robot Control. Several prior JEPAs, such as DINO-WM (Zhou et al., 2024), PLDM (Sobal et al., 2025), and LeWM (Maes et al., 2026b), evaluate control only in simulation, while V-JEPA 2-AC (Assran et al., 2025) plans on physical robots over a frozen pretrained encoder. We evaluate LeWAM on a real robotic platform through direct policy execution, testing whether an end-to-end JEPA can handle real-world contact dynamics and visual variation in closed-loop control.

We evaluate both the end-to-end LeWAM and a variant that supplies pretrained features from a frozen V-JEPA 2.1 (Mur-Labadia et al., 2026) backbone on Pick and Place T, Push-T, and Tape Stack. The physical-robot baselines are Diffusion Policy (Chi et al., 2023) and $\pi_{0.5}$ (Physical Intelligence et al., 2025), a vision-language-action policy. All methods are trained or finetuned on our collected teleoperation data and directly execute the first half of each predicted action chunk during inference. Implementation details appear in Appendix D.

Table 2: Physical robot evaluation. Successful trials out of 30 on Pick and Place T, Push-T, and Tape Stack. GFLOPs count one inference for a full action chunk, including all denoising steps. One full LeWAM action chunk uses 270.2 GFLOPs versus 3603.6 GFLOPs for $\pi_{0.5}$. Figure 1(d) expresses the reciprocal as action chunks per TFLOP, a $13.3\times$ difference in operation-count efficiency.

Method	Pick and Place T	Push-T	Tape Stack	Average	GFLOPs
Diffusion Policy (Chi et al., 2023)	19/30	16/30	0/30	38.9%	270.0
$\pi_{0.5}$ (Physical Intelligence et al., 2025)	14/30	3/30	9/30	28.9%	3603.6
LeWAM (Frozen V-JEPA 2.1 Encoder)	15/30	12/30	5/30	35.6%	270.2
LeWAM	16/30	8/30	6/30	33.3%	270.2

As seen in Table 2, our experiments show that LeWAM is competitive with both baselines, indicating that a JEPA-style world-modeling objective can be trained jointly with action prediction and deployed closed-loop on a *physical robot*. Notably, training the visual encoder end-to-end performs roughly the same as using a pretrained V-JEPA 2.1 encoder frozen across training and inference. This suggests end-to-end training can also learn useful representations for downstream control in a sample-efficient way, despite collapse pressure from the future latent prediction training objective.

Under a reconstruction probe on held-out episodes, both LeWAM variants decode rollouts to comparable visual quality, clearly tracking the manipulated object and robot’s coarse motion while disregarding high-frequency detail (Appendix Figure 16).

5 CONCLUSION AND LIMITATIONS

In this work we introduce LeWAM, an end-to-end JEPA for joint prediction of future latent states and actions, that supports control through both planning and policy execution. On goal-reaching environments, LeWAM demonstrates improved performance in dynamics-based planning over previous JEPA-based world models. On behavioral cloning from expert demonstrations on both simulated and physical robots, LeWAM is competitive with established policy baselines.

Limitations and Future Work.

- **Data Coverage.** Our datasets and task suite are limited in scale, and our complex contact-rich manipulation results use expert demonstrations. The non-expert interaction results on other tasks do not establish that the same training recipe learns difficult contact transitions from suboptimal data. A matched study varying demonstration quality and coverage would test whether broader interaction data improves learned dynamics and physical understanding.
- **Representation Scope.** The reported probes measure physical-state decodability on held-out frames from the same datasets. They do not establish reliable control under distribution shift. Distinctions absent from both recorded actions and their consequences may be compressed; joint supervision does not exclude partial collapse or contextual shortcuts. Testing unfamiliar objects, camera viewpoints, and contact states would clarify this boundary.
- **Single-Task and Single-Embodiment Learning.** Each reported model is trained from scratch for one task and embodiment. These experiments do not test whether the learned action distribution or dynamics transfer across tasks, robots, or instructions. Multitask training across embodiments and conditioning signals (*e.g.*, language) is therefore an open direction.
- **Integration with Hierarchical Planning.** LeWAM provides low-level control and does not include a global high-level planner. The reported horizon sweeps therefore do not establish performance on tasks requiring a long sequence of distinct subgoals. An external planner could decompose such tasks into subgoals, with LeWAM executing each through continuous control.

REPRODUCIBILITY STATEMENT

The experiments in Section 4 use the architectures and baseline configurations in Appendix A, with datasets and evaluation protocols in Appendix B. Anonymous links release all code, configurations, data-processing scripts, result logs, training data, evaluation recordings and results, and redistributable checkpoints.

ETHICS STATEMENT

The experiments collect no new human-subject data. Reused environments, datasets, models, and software remain subject to their original licenses; restricted assets will be accompanied by acquisition instructions. Physical deployment (Section 4.4) requires independent safety mechanisms.

AI USE STATEMENT

LLM-based assistants were primarily used to aid in manuscript polishing and language refinement. This work’s research and manuscript were fully designed, overseen, and executed by its authors. All experimental results and claims were verified manually; the authors take full responsibility for the content of the paper.

REFERENCES

- Mahmoud Assran, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, Yann LeCun, and Nicolas Ballas. Self-supervised learning from images with a joint-embedding predictive architecture. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 15619–15629, 2023. doi: 10.1109/CVPR52729.2023.01499.
- Mido Assran, Adrien Bardes, David Fan, Quentin Garrido, Russell Howes, et al. V-JEPA 2: Self-supervised video models enable understanding, prediction and planning. *arXiv preprint arXiv:2506.09985*, 2025.
- Randall Balestriero and Yann LeCun. LeJEPA: Provable and scalable self-supervised learning without the heuristics. *arXiv preprint arXiv:2511.08544*, 2025.
- Adrien Bardes, Jean Ponce, and Yann LeCun. VICReg: Variance-invariance-covariance regularization for self-supervised learning. In *International Conference on Learning Representations (ICLR)*, 2022.

- Adrien Bardes, Quentin Garrido, Jean Ponce, Xinlei Chen, Michael Rabbat, Yann LeCun, Mahmoud Assran, and Nicolas Ballas. Revisiting feature prediction for learning visual representations from video. *arXiv preprint arXiv:2404.08471*, 2024.
- David Brandfonbrener, Ofir Nachum, and Joan Bruna. Inverse dynamics pretraining learns good representations for multitask imitation. In *Advances in Neural Information Processing Systems*, volume 36, pp. 66953–66978, 2023. doi: 10.52202/075280-2925. URL <https://doi.org/10.52202/075280-2925>.
- Lara Brudermüller, Brandon Hung, Xinghao Zhu, Jiuguang Wang, Nick Hawes, Preston Culbertson, and Simon Le Cleac’h. Generative models from and for sampling-based MPC: A bootstrapped approach for adaptive contact-rich manipulation. *arXiv preprint arXiv:2510.14643*, 2025. URL <https://arxiv.org/abs/2510.14643>.
- Remi Cadene, Simon Aliberts, Francesco Capuano, Michel Aractingi, Adil Zouitine, Pepijn Kooijmans, Jade Choghari, Martino Russi, Caroline Pascal, Steven Palma, Mustafa Shukor, Jess Moss, Alexander Soare, Dana Aubakirova, Quentin Lhoest, Quentin Gallouédec, and Thomas Wolf. LeRobot: An open-source library for end-to-end robot learning. *arXiv preprint arXiv:2602.22818*, 2026.
- Cheng Chi, Siyuan Feng, Yilun Du, Zhenjia Xu, Eric Cousineau, Benjamin C. M. Burchfiel, and Shuran Song. Diffusion policy: Visuomotor policy learning via action diffusion. In *Robotics: Science and Systems (RSS)*, 2023. doi: 10.15607/RSS.2023.XIX.026.
- Thomas W Epps and Lawrence B Pulley. A test for normality based on the empirical characteristic function. *Biometrika*, 70(3):723–726, 1983. doi: 10.1093/biomet/70.3.723.
- Dibya Ghosh, Abhishek Gupta, Ashwin Reddy, Justin Fu, Coline Devin, Benjamin Eysenbach, and Sergey Levine. Learning to reach goals via iterated supervised learning. In *International Conference on Learning Representations (ICLR)*, 2021.
- David Ha and Jürgen Schmidhuber. World models. *arXiv preprint arXiv:1803.10122*, 2018.
- Danijar Hafner, Timothy Lillicrap, Ian Fischer, Ruben Villegas, David Ha, Honglak Lee, and James Davidson. Learning latent dynamics for planning from pixels. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pp. 2555–2565. PMLR, 2019. URL <https://proceedings.mlr.press/v97/hafner19a.html>.
- Danijar Hafner, Timothy Lillicrap, Jimmy Ba, and Mohammad Norouzi. Dream to control: Learning behaviors by latent imagination. In *International Conference on Learning Representations (ICLR)*, 2020.
- Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. Mastering diverse domains through world models. *arXiv preprint arXiv:2301.04104*, 2023.
- Nicklas Hansen, Hao Su, and Xiaolong Wang. TD-MPC2: Scalable, robust world models for continuous control. In *International Conference on Learning Representations (ICLR)*, 2024.
- Nicklas A. Hansen, Hao Su, and Xiaolong Wang. Temporal difference learning for model predictive control. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pp. 8387–8406. PMLR, 2022. URL <https://proceedings.mlr.press/v162/hansen22a.html>.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 770–778, 2016. doi: 10.1109/CVPR.2016.90. URL https://openaccess.thecvf.com/content_cvpr_2016/html/He_Deep_Residual_Learning_CVPR_2016_paper.html.
- Mikael Henaff, William F. Whitney, and Yann LeCun. Model-based planning with discrete and continuous actions. *arXiv preprint arXiv:1705.07177*, 2017.

- Kevin Huang, Sahin Lale, Ugo Rosolia, Yuanyuan Shi, and Anima Anandkumar. CEM-GD: Cross-entropy method with gradient descent planner for model-based reinforcement learning. *arXiv preprint arXiv:2112.07746*, 2021. URL <https://arxiv.org/abs/2112.07746>.
- Petr Ivashkov, Randall Balestriero, and Bernhard Schölkopf. Sensorimotor world models: Perception for action via inverse dynamics. *arXiv preprint arXiv:2606.20104*, 2026.
- Zhenyu Jiang, Yuqi Xie, Kevin Lin, Zhenjia Xu, Weikang Wan, Ajay Mandlekar, Linxi Fan, and Yuke Zhu. DexMimicGen: Automated data generation for bimanual dexterous manipulation via imitation learning. In *IEEE International Conference on Robotics and Automation (ICRA)*, pp. 16923–16930, 2025. doi: 10.1109/ICRA55743.2025.11127809.
- Jyothir S V, Siddhartha Jalagam, Yann LeCun, and Vlad Sobal. Gradient-based planning with world models. *arXiv preprint arXiv:2312.17227*, 2023.
- Yann LeCun. A path towards autonomous machine intelligence. OpenReview, 2022. URL <https://openreview.net/pdf?id=BZ5alr-kVsF>. Position paper, version 0.9.2.
- Shuang Li, Yihuai Gao, Dorsa Sadigh, and Shuran Song. Unified video action model. *arXiv preprint arXiv:2503.00200*, 2025.
- Weixin Liang, Lili Yu, Liang Luo, Srinivasan Iyer, Ning Dong, Chunting Zhou, Gargi Ghosh, Mike Lewis, Wen-tau Yih, Luke Zettlemoyer, and Xi Victoria Lin. Mixture-of-Transformers: A sparse and scalable architecture for multi-modal foundation models. *Transactions on Machine Learning Research*, 2025. URL <https://arxiv.org/abs/2411.04996>.
- Yaron Lipman, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching for generative modeling. In *International Conference on Learning Representations (ICLR)*, 2023.
- Lucas Maes, Quentin Le Lidec, Luiz Facury, Nassim Massaudi, Ayush Chaurasia, Francesco Capuano, Richard Gao, Taj Gillin, Dan Haramati, Damien Scieur, Yann LeCun, and Randall Balestriero. stable-worldmodel: A platform for reproducible world modeling research and evaluation. *arXiv preprint arXiv:2605.21800*, 2026a.
- Lucas Maes, Quentin Le Lidec, Damien Scieur, Yann LeCun, and Randall Balestriero. LeWorldModel: Stable end-to-end joint-embedding predictive architecture from pixels. *arXiv preprint arXiv:2603.19312*, 2026b.
- Ajay Mandlekar, Danfei Xu, Josiah Wong, Soroush Nasiriany, Chen Wang, Rohun Kulkarni, Li Fei-Fei, Silvio Savarese, Yuke Zhu, and Roberto Martín-Martín. What matters in learning from offline human demonstrations for robot manipulation. In *Proceedings of the 5th Conference on Robot Learning*, volume 164 of *Proceedings of Machine Learning Research*, pp. 1678–1690. PMLR, 2022. URL <https://proceedings.mlr.press/v164/mandlekar22a.html>.
- Vincent Micheli, Eloi Alonso, and François Fleuret. Transformers are sample-efficient world models. In *International Conference on Learning Representations (ICLR)*, 2023.
- Lorenzo Mur-Labadia, Matthew Muckley, Amir Bar, Mido Assran, Koustuv Sinha, Mike Rabbat, Yann LeCun, Nicolas Ballas, and Adrien Bardes. V-JEPA 2.1: Unlocking dense features in video self-supervised learning. *arXiv preprint arXiv:2603.14482*, 2026.
- Hoang Nguyen, Xiaohao Xu, and Xiaonan Huang. Latent geometry beyond search: Amortizing planning in world models. *arXiv preprint arXiv:2605.08732*, 2026.
- NVIDIA. Cosmos 3: Omnimodal world models for physical AI. *arXiv preprint arXiv:2606.02800*, 2026.
- Maxime Oquab, Timothée Darcet, Théo Moutakanni, et al. DINOv2: Learning robust visual features without supervision. *arXiv preprint arXiv:2304.07193*, 2023.
- Tao Pang, H. J. Terry Suh, Lujie Yang, and Russ Tedrake. Global planning for contact-rich manipulation via local smoothing of quasi-dynamic contact models. *IEEE Transactions on Robotics*, 39(6):4691–4711, 2023. doi: 10.1109/TRO.2023.3300230.

- Seohong Park, Kevin Frans, Benjamin Eysenbach, and Sergey Levine. OGBench: Benchmarking offline goal-conditioned RL. In *International Conference on Learning Representations (ICLR)*, 2025.
- William Peebles and Saining Xie. Scalable diffusion models with transformers. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, pp. 4172–4182, 2023. doi: 10.1109/ICCV51070.2023.00387.
- Physical Intelligence, Kevin Black, et al. $\pi_{0.5}$: A vision-language-action model with open-world generalization. *arXiv preprint arXiv:2504.16054*, 2025.
- Cristina Pinneri, Shambhuraj Sawant, Sebastian Blaes, Jan Achterhold, Joerg Stueckler, Michal Rolinek, and Georg Martius. Sample-efficient cross-entropy method for real-time planning. In *Proceedings of the 2020 Conference on Robot Learning*, volume 155 of *Proceedings of Machine Learning Research*, pp. 1049–1065. PMLR, 2021. URL <https://proceedings.mlr.press/v155/pinneri21a.html>.
- Olivier Roy and Martin Vetterli. The effective rank: A measure of effective dimensionality. In *15th European Signal Processing Conference (EUSIPCO)*, pp. 606–610, 2007. URL <https://www.eurasip.org/Proceedings/Eusipco/Eusipco2007/Papers/a5p-h05.pdf>.
- Reuven Rubinstein. The cross-entropy method for combinatorial and continuous optimization. *Methodology and Computing in Applied Probability*, 1(2):127–190, 1999. doi: 10.1023/A:1010091220143.
- Vlad Sobal, Jyothir S V, Siddhartha Jalagam, Nicolas Carion, Kyunghyun Cho, and Yann LeCun. Joint embedding predictive architectures focus on slow features. In *NeurIPS Workshop on Self-Supervised Learning: Theory and Practice*, 2022. URL <https://arxiv.org/abs/2211.10831>.
- Vlad Sobal, Wancong Zhang, Kyunghyun Cho, Randall Balestriero, Tim G. J. Rudner, and Yann LeCun. Learning from reward-free offline data: A case for planning with latent dynamics models. *arXiv preprint arXiv:2502.14819*, 2025.
- Yuhai Wang, Jiawei Xia, Rongxuan Zhou, Xiao Hu, Yongliang Shi, Jing Du, and Yang Ye. PRISM: PRior-guided Imagination Sampling in world Models. *arXiv preprint arXiv:2606.07974*, 2026.
- Grady Williams, Nolan Wagener, Brian Goldfain, Paul Drews, James M. Rehg, Byron Boots, and Evangelos A. Theodorou. Information theoretic MPC for model-based reinforcement learning. In *IEEE International Conference on Robotics and Automation (ICRA)*, pp. 1714–1721, 2017. doi: 10.1109/ICRA.2017.7989202.
- Haiyu Wu, Randall Balestriero, and Morgan Levine. VISReg: Variance-invariance-sketching regularization for JEPA training. *arXiv preprint arXiv:2606.02572*, 2026.
- Jun Yamada, Shaohong Zhong, Jack Collins, and Ingmar Posner. D-Cubed: Latent diffusion trajectory optimisation for dexterous deformable manipulation. *arXiv preprint arXiv:2403.12861*, 2024. URL <https://arxiv.org/abs/2403.12861>.
- Wilson Yan, Yunzhi Zhang, Pieter Abbeel, and Aravind Srinivas. Videogpt: Video generation using vq-vae and transformers. *arXiv preprint arXiv:2104.10157*, 2021.
- Seonghyeon Ye, Yunhao Ge, Kaiyuan Zheng, Shenyuan Gao, Sihyun Yu, et al. World action models are zero-shot policies. *arXiv preprint arXiv:2602.15922*, 2026.
- Tianyuan Yuan, Zibin Dong, Yicheng Liu, and Hang Zhao. Fast-WAM: Do world action models need test-time future imagination? *arXiv preprint arXiv:2603.16666*, 2026.
- Gaoyue Zhou, Hengkai Pan, Yann LeCun, and Lerrel Pinto. DINO-WM: World models on pre-trained visual features enable zero-shot planning. *arXiv preprint arXiv:2411.04983*, 2024.
- Chuning Zhu, Raymond Yu, Siyuan Feng, Benjamin Burchfiel, Paarth Shah, and Abhishek Gupta. Unified world models: Coupling video and action diffusion for pretraining on large robotic datasets. *arXiv preprint arXiv:2504.02792*, 2025.

Appendix

Supplementary material for **LeWAM: End-to-End World and Action Modeling with JEPAs**

Contents

A	Implementation Details	p. 15
	LeWAM Planning	p. 17
	Baseline Implementations	p. 17
B	Benchmarks and Evaluation	p. 18
	Tasks and Datasets	p. 18
	Goals and Success Criteria	p. 19
	Planning and Timing Protocols	p. 21
C	Additional Results	p. 21
	Visual Planning Results	p. 21
	Component and Visual-Encoder Ablations	p. 22
	Probing Latent Representations	p. 23
	Visual Reconstructions and Open-Loop Rollouts	p. 24
	Multimodal Action Distributions	p. 27
	Simulation Task-Completion Results	p. 27
D	Physical-Robot Setup and Evaluation	p. 27
	Platform and Data Collection	p. 27
	Model Implementation	p. 28
	Physical-Robot Baseline Configurations	p. 28
	Closed-Loop Evaluation	p. 29
	Representation Diagnostics and Evaluation	p. 29
E	Discussion of the Role of Joint Prediction	p. 31
	Representation and Dynamics	p. 31
	Planning with Learned Action Proposals	p. 32
F	Author Contribution Statement	p. 33

A IMPLEMENTATION DETAILS

Temporal Sampling. Given observation-history stride Δ , the predictor in Section 3.1 uses two history latents $\mathbf{z}_{\leq t} = (\mathbf{z}_{t-\Delta}, \mathbf{z}_t)$, a dynamics stride of $K = 5$ raw steps, and an action chunk of $C = 2K = 10$ raw steps. We use a stride of $\Delta = K$ for goal-reaching tasks and $\Delta = 1$ for task-completion frames. To predict $\mathbf{z}_{t+5}$, the dynamics branch receives all five actions $\mathbf{a}_t, \dots, \mathbf{a}_{t+4}$; the flow policy models all ten actions $\mathbf{a}_t, \dots, \mathbf{a}_{t+9}$. Only frames are subsampled. During imagined rollouts, the history update retains the two most recent latent states.

Inner Architecture. Unless a task-specific deviation is stated, the observation encoder is a ResNet-18-style network (He et al., 2016) trained from scratch. Its convolutional backbone uses GroupNorm with 16 groups. Input images are cropped from 224×224 to 202×202 , with random offsets during training and a center crop at evaluation. A 1×1 convolution produces 32 spatial maps; spatial softmax pooling converts these to 32 expected normalized image-coordinate pairs. A two-layer projection MLP with BatchNorm maps these 64 features through a hidden width of 384 to a 192-dimensional latent state. The MoT predictor has width 192, four blocks, four attention heads, and dropout 0.1. Its token sequence contains two history latents, one learned state query, five clean action tokens, and ten action flow slots. Learned type and positional embeddings distinguish history, state query, clean action, and action flow tokens. The state readout uses LayerNorm followed by a linear map to the 192-dimensional latent, while the action readout maps each action flow slot to a velocity vector in raw-action space.

Goal information $\mathbf{c}_t$ is optional and is omitted for goal-free task completion. For goal-conditioned experiments, $\mathbf{c}_t$ contains the goal latent $\mathbf{z}_g$ and the normalized remaining horizon. The goal latent is supplied as an additional token in the action branch; only noisy action tokens can attend to it. The horizon h between the current timestep and the goal is normalized into a $[0, 1]$ range relative to a maximal value $H_{\max}$ which dictates the longest distance-from-goal seen during training. h modulates the noisy action tokens using AdaLN-Zero (Peebles & Xie, 2023) blocks with hidden width 512 and a 128-dimensional sinusoidal embedding.

In total, the model has approximately 16.9M trainable parameters. For the simulated experiments, action sampling uses eight Euler steps to balance speed and quality. The reported MSE dynamics configuration uses an online encoder target, no state noise, no state flow integration, no momentum target encoder, only SIGReg with 1,024 projections and 17 knots. For the training objective in Equation 5, we set $\lambda_{\text{sig}} = 0.04$.

Training Recipe in Simulated Environments. For goal-conditioned visual planning, LeWAM runs use a batch size of 128 and train for 50 epochs, with a base learning rate of 1.5×10^{-4} and five warmup epochs. Notably, in most environments, training for about 10 epochs can converge. The goal-free task-completion runs use a batch size of 64 and train for 120 epochs, with a base learning rate of 10^{-4} and ten warmup epochs. Both use AdamW with weight decay 10^{-4} and a uniform learning rate across the encoder and MoT predictor. The learning rate increases linearly during warmup and then follows an epochwise cosine schedule. Training uses full precision, a gradient-norm clip of 1, and seed 42. We split valid decision points into 90% training and 10% validation, then sample training points with replacement within each epoch. After each epoch, we retain the checkpoint with the lowest validation action flow-matching loss, where task success is not used for checkpoint selection.

Tokens and Conditioning. With the optional goal context $\mathbf{c}_t$ omitted, the MoT sequence is

$$\left[\mathbf{z}_{t-\Delta}, \mathbf{z}_t, \mathbf{q}_{t+K}^z, \mathbf{A}_{t,K}, \mathbf{A}_{t,C}^{(s)} \right]. \quad (9)$$

Here $\mathbf{A}_{t,K}$ and $\mathbf{A}_{t,C}^{(s)}$ expand into K clean and $2K$ noisy action tokens, respectively. The first three tokens are processed by f_z , and both action groups by f_a . Each MoT branch has separate query, key, value, output, normalization, and feedforward parameters, while one global masked self-attention operates over the complete token sequence. Flow timestep s modulates only the noisy action slots through AdaLN-Zero.




Figure 10: **Token-Level MoT Attention.** Expansion of Figure 3 for two history latents, one learned state query, five clean actions, and ten interpolated action tokens. Rows are queries; columns are keys. Filled cells permit attention; white cells are masked. Blue and brown mark latent state and action velocity prediction queries; gray marks context. Wider gaps separate token groups. Optional goal conditioning c_t is omitted. Δ is the history stride, and action indices remain at the raw environment rate.

Attention Mask. Let $\mathcal{V}(x)$ denote the tokens visible to query x . The mask below describes the sequence without optional goal context c_t . When this context is provided, the goal token is additionally visible to the noisy action tokens only. Action prefixes begin at the current chunk’s start t ; padded history keys are excluded:

$$\begin{aligned}
 \mathcal{V}(z_i) &= \{z_{t-\Delta}, z_t\}, & i &\in \{t-\Delta, t\}, \\
 \mathcal{V}(q_{t+K}^z) &= \{z_{t-\Delta}, z_t, q_{t+K}^z, \mathbf{A}_{t,K}\}, \\
 \mathcal{V}(a_j) &= \{z_{t-\Delta}, z_t, \mathbf{a}_{\leq j}\}, & j &\in \{t, \dots, t+K-1\}, \\
 \mathcal{V}(a_j^{(s)}) &= \{z_{t-\Delta}, z_t, \mathbf{a}_{\leq j}^{(s)}\}, & j &\in \{t, \dots, t+2K-1\}.
 \end{aligned} \tag{10}$$

Clean action tokens therefore cannot relay target information to history or action flow tokens across layers. A visualization of the attention mask is provided in Figure 10.

A.1 LEWAM PLANNING

Optimization Details. The goal cost $\mathcal{C}$ is defined in Eq. (7). During refinement, model parameters remain frozen and only the normalized action plans receive gradients. For each sampled plan $\mathbf{A}^{\text{sample}}$, Adam minimizes the goal cost with a quadratic proximity penalty:

$$\mathcal{J}_{\text{prox}}(\mathbf{A}; \mathbf{A}^{\text{sample}}) = \mathcal{C}(\mathbf{A}) + \lambda_{\text{prox}} \|\mathbf{A} - \mathbf{A}^{\text{sample}}\|_F^2. \quad (11)$$

The penalty affects the updates, while candidate retention and final selection use the unpenalized goal cost $\mathcal{C}$. We clip normalized actions to $[-a_{\max}, a_{\max}]$ after each update.

Algorithm 1 Candidate refinement and selection for LeWAM planning

Require: $\mathbf{z}_{\leq t}, \mathbf{c}_t, \mathbf{z}_g, L, N, J, \eta, \lambda_{\text{prox}}, a_{\max}$
Sample candidates from the learned policy
1: $\mathbf{A}_{1:N}^{\text{sample}} \leftarrow N$ policy rollouts from $(\mathbf{z}_{\leq t}, \mathbf{c}_t)$ for L/K blocks
2: $c_{1:N} \leftarrow \mathcal{C}(\mathbf{A}_{1:N}^{\text{sample}})$
3: $\mathbf{A}_{1:N}^* \leftarrow \mathbf{A}_{1:N}^{\text{sample}}, \quad \mathbf{A}_{1:N} \leftarrow \mathbf{A}_{1:N}^{\text{sample}}$
4: $c_{1:N}^* \leftarrow c_{1:N}$
Refine each candidate
5: **for** $j = 0, \dots, J - 1$ **do**
6: $\mathbf{A}_{1:N} \leftarrow \text{AdamStep}(\mathbf{A}_{1:N}, \nabla_{\mathbf{A}_{1:N}} \mathcal{J}_{\text{prox}}(\mathbf{A}_{1:N}; \mathbf{A}_{1:N}^{\text{sample}}), \eta)$
7: $\mathbf{A}_{1:N} \leftarrow \text{clip}(\mathbf{A}_{1:N}, -a_{\max}, a_{\max})$
8: $c_{1:N} \leftarrow \mathcal{C}(\mathbf{A}_{1:N})$
9: $\mathbf{A}_n^* \leftarrow \mathbf{A}_n, \quad c_n^* \leftarrow c_n \quad \forall n \text{ such that } c_n < c_n^*$
10: **end for**
Select the lowest-cost iterate
11: $n^* \leftarrow \arg \min_{n=1, \dots, N} c_n^*$
12: **return** $\mathbf{A}_{n^*}^*$

Unless specified otherwise, Adam uses learning rate 0.05, gradient-norm clipping at 10, $a_{\max} = 3$, and $\lambda_{\text{prox}} = 0.01$.

A.2 BASELINE IMPLEMENTATIONS

The baselines in Section 4.1 use the released Stable World Model codebase (Maes et al., 2026a), retaining its environment wrappers, dataset interface, action bounds, and evaluation loop. Unless stated below, each method follows its official implementation and hyperparameters.¹ These configurations describe goal-conditioned comparisons. Methods share observations, actions, start and goal pairs, and success predicates. We list deviations from released configurations.

- **LeWM (IDM).** This is the GC-IDM (Nguyen et al., 2026) configuration: a pretrained LeWM checkpoint trained by us is frozen, then a goal-conditioned inverse-dynamics model is post-trained. We follow its official two-stage implementation. Goal reaching samples a remaining horizon uniformly from 1 to 50 and clips it to the episode boundary.
- **LeWM (CEM).** We follow the official LeWM training and CEM evaluation settings (Maes et al., 2026b). On the four inherited task types, our reproduced checkpoints perform similarly to the released checkpoints but were unable to match the reported scores. All quantitative experiments using LeWM as a baseline use checkpoints trained by us. The `Reacher` data differs from the released random-policy collection, as detailed in Appendix B.1.
- **DINO-WM.** We evaluate the released patch-token model and a matched CLS-token variant; tokenization is their only architectural difference. The six added tasks use ten training epochs, while the original four retain their four-epoch checkpoints.
- **GCBC.** We use the goal-conditioned behavior-cloning baseline from the released codebase (Maes et al., 2026a), with its pretrained DINOv2 encoder frozen and the policy trained for 20 rather than

¹Official repositories: DINO-WM, Stable World Model, and GC-IDM.

100 epochs. Short-horizon goal reaching samples a uniformly later frame from the same episode. At evaluation, the policy executes its predicted five-step action chunks without planning.

Behavior-Cloning Baseline in Simulation. The task-completion comparison in Table 10 uses a policy trained on demonstrations from each task, conditioned on observation history without a goal image or goal-horizon input.

- **Diffusion Policy** (Chi et al., 2023). We use the convolutional variant (DP-C), with a ResNet-18 encoder, GroupNorm, and spatial softmax for each camera. The image features condition a one-dimensional U-Net with channel widths (512, 1024, 2048), kernel size five, and diffusion-step embedding dimension 128. Training minimizes noise-prediction MSE with a 100-step squared-cosine diffusion schedule; inference uses 100 DDPM steps. The policy takes two consecutive observations, predicts a 16-action sequence, and executes eight actions starting at the current timestep before resampling.

ToolHang uses the scene camera; DrawerCleanup and Transport use the scene camera and both wrist cameras. Inputs are 224×224 RGB images, with random 202×202 crops during training and center crops at evaluation, normalized to $[-1, 1]$. The policy uses visual observations only. Training windows advance at the raw frame rate, with boundary frames and actions repeated for padding. We randomly split windows into 90% training and 10% validation using seed 42. Each policy trains for 120 epochs with batch size 64 and AdamW, learning rate 10^{-4} , betas (0.95, 0.999), and weight decay 10^{-6} . The learning rate follows cosine decay after 500 warmup updates. Actions retain their recorded scale for ToolHang and Transport; DrawerCleanup uses min-max scaling to $[-1, 1]$. We select the checkpoint with the lowest validation denoising loss and evaluate its raw weights without exponential moving averaging. Each checkpoint is evaluated with seeds 0, 1, and 42, using 50 episodes per seed.

B BENCHMARKS AND EVALUATION

B.1 TASKS AND DATASETS

Task Selection. We retain the four environments used by LeWM (Maes et al., 2026b) for direct comparison, then add six tasks that broaden navigation and manipulation to larger mazes, interacting objects, coupled button states, precise tool placement, bimanual transfer, and articulated drawers (Figure 5). The ten tasks also span different action dimensions, observation views, trajectory lengths, and data-collection regimes. Together, they test the same model design under varied control and data conditions.

Data Provenance. We describe each dataset in the environment order of Figure 5.

- **Two-Room.** We use the dataset released with LeWM (Maes et al., 2026b), collected by a noisy heuristic navigation policy.
- **Reacher.** We collect goal-directed trajectories with a noisy policy in a visible-target environment. These are distinct from the independent-uniform-action data collected by the released Stable World Model (Maes et al., 2026a) collector, whose trajectories average 201 frames. We keep the random, goal-directed, and equally mixed collections distinct; Table 3 reports the goal-directed data.
- **Push-T.** We use the non-expert interaction dataset released by DINO-WM (Zhou et al., 2024), generated by replaying the original expert trajectories with different levels of action noise.
- **Cube.** We use the dataset released with LeWM (Maes et al., 2026b) for the single-cube OGBench (Park et al., 2025) environment, collected by a scripted expert that picks and places the cube at sampled target locations.
- **PointMaze-Large.** We use the official OGBench (Park et al., 2025) pointmaze-large-navigate-v0 dataset, collected by a noisy expert that navigates between randomly sampled locations. We render its recorded simulator states into RGB observations.

- **Scene.** We collect data ourselves using the Stable World Model (Maes et al., 2026a) wrapper for OGBench (Park et al., 2025) Scene (swm/OGBScene-v0) and its ExpertPolicy, a noisy, closed-loop scripted oracle. The policy interacts with the cube, drawer, window, and buttons; we record a single camera view.
- **Puzzle.** We collect data ourselves using the Stable World Model (Maes et al., 2026a) wrapper for the 3×3 OGBench (Park et al., 2025) Puzzle environment and its scripted ExpertPolicy. The policy presses buttons in sampled sequences; we record a single camera view.
- **ToolHang.** We use the official human behavior-cloning demonstrations from robomimic (Mandlekar et al., 2022).
- **Transport.** We use the official generated behavior-cloning demonstrations from DexMimicGen (Jiang et al., 2025).
- **DrawerCleanup.** We use the official generated behavior-cloning demonstrations from DexMimicGen (Jiang et al., 2025).

The three behavior-cloning datasets broaden the visual-planning comparison to complex contact-rich behaviors with higher-dimensional actions, although their task-focused demonstrations may limit the diversity of sampled visual goals.

Table 3 summarizes all ten tasks, including observation views, raw action dimensions, and dataset scale. Eight rows have audited source counts; Scene and Puzzle use author-reported statistics. Counts for the random and mixed Reacher variants are not included. Recorded lengths describe the data, while Appendix B.2 specifies evaluation horizons and execution budgets. Recorded duration uses the native sampling rate before temporal subsampling; synchronized views of the same trajectories are counted once. ToolHang uses one oblique scene view (about 45°), while Transport and DrawerCleanup use a scene view and both wrist views.

Table 3: Simulation Benchmark Datasets. Ten tasks in Figure 5 order. Views gives camera-stream count; Dim. is the raw action dimension. Statistics cover complete source datasets before splits. Eight rows are audited; Scene and Puzzle are author reported. Total hours are computed before rounding per-task values.

Task	Collection regime	Views	Dim.	Traj.	Mean frames	Total frames	Hz	Hours
Two-Room	Noisy heuristic	1	2	10,000	92	920,809	10	25.58
Reacher	Goal-directed policy	1	2	17,504	143	2,498,208	25	27.76
Push-T	Non-expert	1	2	18,685	125	2,336,736	10	64.91
Cube	Scripted expert	1	5	10,000	201	2,010,000	20	27.92
PointMaze-Large	Scripted navigation	1	2	1,000	1,000	1,000,000	10	27.78
Scene	Scripted expert	1	5	10,024	200	2,004,800	20	27.84
Puzzle (3x3)	Scripted expert	1	5	10,024	200	2,004,800	20	27.84
ToolHang	Human demos	1	7	200	480	95,962	20	1.33
Transport	Generated demos	3	14	1,029	408	419,712	20	5.83
DrawerCleanup	Generated demos	3	24	1,026	291	298,235	20	4.14
Reported total (10)				79,492		13,589,262		240.93

B.2 GOALS AND SUCCESS CRITERIA

For the comparisons in Section 4.1, unless stated otherwise, we report mean success and standard deviation over three seeds with 50 episodes per seed. Methods within a protocol share observations, action limits, execution budgets, and success predicates.

Goal-Conditioned Visual Planning. In Section 4.2, the initial state is sampled from an offline trajectory and the frame 25 steps later supplies the visual goal. Each controller receives 50 environment actions across all ten environments. We report first-hit success for comparison with the released LeWM evaluation (Maes et al., 2026b); the Cube predicate correction is detailed below. Scene and Puzzle follow the same protocol; their random-policy and zero-action controls appear alongside those for the other environments in Table 4.

Environment	Raw action dim.	25-step goal reaching		Task Completion	
		Zero-action	Random	Zero-action	Random
Two-Room	2	7.3 ± 0.9	25.3 ± 5.0	0.0 ± 0.0	2.7 ± 2.5
Reacher	2	3.3 ± 1.9	12.0 ± 1.6	0.0 ± 0.0	12.0 ± 4.9
Push-T	2	0.0 ± 0.0	1.3 ± 0.9	0.0 ± 0.0	0.0 ± 0.0
Cube	5	42.7 ± 6.6	44.0 ± 1.6	4.7 ± 2.5	6.0 ± 1.6
Cube (+EEF)	5	2.7 ± 0.9	4.7 ± 0.9	4.7 ± 2.5	6.0 ± 1.6
PointMaze	2	6.0 ± 0.0	9.3 ± 1.9	0.7 ± 0.9	6.0 ± 3.3
Scene	5	2.7 ± 2.3	2.7 ± 2.3	0.0 ± 0.0	0.0 ± 0.0
Puzzle	5	0.7 ± 1.2	0.7 ± 1.2	0.0 ± 0.0	0.0 ± 0.0
ToolHang	7	22.7 ± 2.5	26.0 ± 1.6	0.0 ± 0.0	0.0 ± 0.0
Transport	14	3.3 ± 0.9	2.7 ± 0.9	0.0 ± 0.0	0.0 ± 0.0
DrawerCleanup	24	0.0 ± 0.0	6.0 ± 0.0	0.0 ± 0.0	0.0 ± 0.0

Darker shading indicates a higher control success rate.

Table 4: **Action Dimensions and Policy-Floor Controls.** Rows cover all ten environments in Figure 5, plus a Cube (+EEF) row isolating the corrected end-effector predicate. Raw action dimensionality is reported beside zero-action and random-policy success rates (%) for short-horizon goal reaching and full-trajectory task completion, averaged over three seeds. These are control floors, not additional learned-policy experiments. Cube has the same task-completion values in both rows because the end-effector correction applies only to short-horizon goal reaching. DrawerCleanup exposes 24 raw coordinates; the policy predicts 22 active dimensions after removing two constant joints. PointMaze uses the PointMaze-Large layout.

Policy-Floor Controls. Table 4 also reports zero-action and random-policy success under the full-trajectory task-completion evaluation. These blind controls cover the benchmark environments; the learned task-completion policies in Section 4.4 are compared on DrawerCleanup, Transport, and ToolHang only.

Cube Success Predicate. Specifically, we add an end-effector (EEF) check to the Cube success criterion to exclude spurious successes. The original Cube criterion checks whether the cube center lies within 0.04 m of the target, regardless of the gripper pose. For 25-step goals, this credits zero action on 42.7% of episodes (Table 4). The corrected criterion additionally requires the end effector to lie within 0.04 m of its goal-frame position, reducing this floor to 2.7%. Figure 11 illustrates the object-only failure case.



Figure 11: **A Degenerate Cube Success under Short-Horizon Goal Reaching.** Each row shows a policy rollout (left, frames over time) beside its goal frame (far right, after the divider). Throughout the execution the gripper never reaches the goal end-effector pose, yet the cube stays within the 0.04 m object-only threshold of the goal, so OGBench’s object-position-only predicate scores the episode a success. Our end-effector term (Table 4) rejects these degenerate successes by additionally requiring the gripper to reach the goal pose.

B.3 PLANNING AND TIMING PROTOCOLS

Goal-Horizon Sweep. Figure 7 in Section 4.2 and Table 5 in Appendix C.1 use the following protocols. On `Push-T`, LeWAM uses SIGReg and the sampling and gradient-refinement procedure in Section 3.2; LeWM (CEM) (Maes et al., 2026b) uses a checkpoint trained by us with the reported 300×30 CEM configuration. For `Push-T`, the goal is the recorded frame 25, 50, 75, or 100 steps after the initial frame. Each episode starts at frame 0, receives twice the goal horizon in planning budget (total actions executed), and replans every 25 steps. These are the reported `Push-T` solver configurations; matched-compute comparisons remain pending. `Cube` uses a separately reported horizon sweep with five gradient steps and a fixed starting point. Its 25-step LeWM (CEM) result differs from the sampled-start evaluation in Figure 6, so the two protocols should not be conflated. For `Scene`, the author-reported sweep gives four goal offsets; its start is the first trajectory frame, while its action budget and replanning interval are not specified in the supplied result summary. For `Puzzle`, both reported sweeps cover the same four goal offsets from the first trajectory frame; LeWM (CEM) uses the same baseline settings as the other evaluations.

Planner-Sensitivity Protocol. On `Push-T`, we vary Adam refinement updates while fixing the candidate count at 32, then vary the number of sampled candidates while fixing refinement at five updates. Goal-reaching success and per-plan latency are measured over three seeds with 50 episodes per seed. The measured curves appear in Figure 12 in Appendix C.1.

Timing Scope. Table 6 in Appendix C.1 reports time per planner call. Figure 1(c) averages the reciprocal throughput of each of its four tasks, using 1000/milliseconds per plan to obtain plans/s. The compared settings are 300×30 CEM for LeWM (Maes et al., 2026b) and five gradient steps for LeWAM. Hardware and the exact timing boundaries remain to be documented.

C ADDITIONAL RESULTS

C.1 VISUAL PLANNING RESULTS

Table 5 gives the measured goal-horizon sweeps behind Figure 7. All four `Cube` horizon points have reported means and standard deviations. Table 6 gives the per-call timing behind Figure 1(c). Appendix B.3 specifies the evaluation and timing protocols.

Table 5: **Goal-Reaching Success across Horizons.** Success rate (%), mean $\pm$ standard deviation. Measured values for Figure 7. The `Cube` horizon sweep is distinct from the sampled-start 25-step evaluation in Figure 6.

Task	Method	Goal horizon (steps)			
		25	50	75	100
<code>Push-T</code>	LeWAM	94.7 $\pm$ 3.1	60.7 $\pm$ 3.1	44.0 $\pm$ 3.5	43.3 $\pm$ 4.6
	LeWM (CEM)	68.7 $\pm$ 3.1	12.0 $\pm$ 4.0	3.3 $\pm$ 2.3	2.7 $\pm$ 1.2
<code>Cube</code>	LeWAM	100.0 $\pm$ 0.0	88.0 $\pm$ 2.0	80.7 $\pm$ 2.3	82.0 $\pm$ 2.0
	LeWM (CEM)	100.0 $\pm$ 0.0	0.0 $\pm$ 0.0	4.0 $\pm$ 3.5	4.0 $\pm$ 3.5
<code>Scene</code>	LeWAM	82.7 $\pm$ 3.1	73.3 $\pm$ 8.1	65.3 $\pm$ 3.1	56.0 $\pm$ 4.0
	LeWM (CEM)	44.7 $\pm$ 13.0	31.3 $\pm$ 2.3	37.3 $\pm$ 5.8	36.0 $\pm$ 6.0
<code>Puzzle</code>	LeWAM	74.0 $\pm$ 12.5	14.0 $\pm$ 3.5	4.0 $\pm$ 3.5	3.3 $\pm$ 1.2
	LeWM (CEM)	10.0 $\pm$ 0.0	4.7 $\pm$ 2.3	4.0 $\pm$ 2.0	0.7 $\pm$ 1.2

Table 6: **Planning Speed across Four Visual-Goal Tasks.** Each plan is one planner call. LeWM (CEM) (Maes et al., 2026b) uses 300×30 CEM; LeWAM uses five gradient steps. Action dimensions are raw environment action dimensions. Throughput is the reciprocal of time per plan. The mean row averages per-task throughput, and its ratio compares those two means.

Task	Action dim.	LeWM (CEM)		LeWAM		Speedup
		ms/plan	plans/s	ms/plan	plans/s	
Scene	5	705	1.42	18.8	53.19	37.5×
Puzzle	5	731	1.37	16.3	61.35	44.8×
Cube	5	771	1.30	21.1	47.39	36.5×
Push-T	2	414	2.42	20.8	48.08	19.9×
Task mean	—	—	1.62	—	52.50	32.3×

Planner Sensitivity. In the `Push-T` sweep in Figure 12, success plateaus after five refinement updates while latency continues to rise. Increasing the number of sampled candidates beyond eight gives no consistent success gain in this sweep.




Figure 12: **Planner Sensitivity on Push-T.** Goal-reaching success (mean and standard deviation across three seeds, 50 episodes each) and per-plan latency versus (left) Adam refinement updates with 32 candidate plans and (right) sampled candidate plans with five refinement updates.

C.2 COMPONENT AND VISUAL-ENCODER ABLATIONS

Table 7 records the SIGReg and action-head comparisons in Figure 8(a,c). All variants use epoch-30 checkpoints. Table 8 reports the separate encoder comparison in Figure 8(b). Direct control executes the goal-conditioned policy without search. Planning performs test-time action search through the learned dynamics. The encoder panel reports equal-weight averages over `Push-T`, `Cube`, `Scene`, and `Puzzle`.

Table 7: **SIGReg and Action-Head Ablations.** Direct and planning success rates (%) using epoch-30 checkpoints for all variants. Gain is planning minus direct success in percentage points. Standard deviations are reported where available; none is inferred for gains.

Environment	Variant	Direct	Planning	Gain (pp)
Push-T	LeWAM	82.7 ± 8.1	97.3 ± 3.1	+14.7
	w/o SIGReg	69.3 ± 10.1	72.0 ± 3.5	+2.7
	MSE action head	90.0 ± 4.0	95.3 ± 1.2	+5.3
Scene	LeWAM	92.0 ± 5.3	97.3 ± 1.2	+5.3
	w/o SIGReg	91.3 ± 2.3	93.3 ± 3.1	+2.0
	MSE action head	89.3 ± 4.2	90.0 ± 2.0	+0.7
Puzzle	LeWAM	76.0 ± 7.2	80.7 ± 11.0	+4.7
	w/o SIGReg	67.3 ± 9.5	75.3 ± 9.0	+8.0
	MSE action head	76.7 ± 6.1	74.7 ± 6.1	-2.0
Cube	LeWAM	99.3 ± 1.2	100.0	+0.7
	w/o SIGReg	99.3 ± 1.2	99.3 ± 1.2	0.0
	MSE action head	98.0 ± 2.0	99.3 ± 1.2	+1.3

Table 8: **Visual Encoders for Direct Control and Planning.** Success rate (%) with the same goal-conditioned evaluation protocol for all encoder variants. Direct executes the policy without search; planning uses test-time action search. Gain is planning minus direct success in percentage points. Averages are equal-weight means over the four tasks; gains use unrounded means.

Visual encoder	Mode	Push-T	Cube	Scene	Puzzle	Average
LeWAM (end-to-end)	Direct	85.3 ± 5.0	100.0 ± 0.0	88.0 ± 6.9	73.3 ± 5.8	86.7
	Planning	98.7 ± 1.2	100.0 ± 0.0	93.3 ± 3.1	74.0 ± 12.5	91.5
	Gain (pp)	+13.4	0.0	+5.3	+0.7	+4.9
frozen DINOv2 CLS	Direct	72.7 ± 7.6	98.7 ± 1.2	82.0 ± 4.0	62.7 ± 7.0	79.0
	Planning	78.7 ± 6.4	99.3 ± 1.2	91.3 ± 1.2	61.3 ± 13.0	82.7
	Gain (pp)	+6.0	+0.6	+9.3	-1.4	+3.6
frozen DINOv2 patch mean	Direct	66.0 ± 12.5	93.3 ± 2.3	84.0 ± 2.0	60.0 ± 7.2	75.8
	Planning	67.3 ± 8.1	98.7 ± 1.2	86.0 ± 4.0	58.7 ± 7.6	77.7
	Gain (pp)	+1.3	+5.4	+2.0	-1.3	+1.9

C.3 PROBING LATENT REPRESENTATIONS

Table 9 extends the probing results presented on Push-T in Section 4.3 to Cube, Scene, Puzzle. All probes report R^2 from ridge regression on a 4000/1000-frame train/test split. *Object* is the mean R^2 over object ground-truth variables and *Proprio* over proprioceptive variables. *Action (IDM)* predicts the action chunk from z_t (and from $[z_t, z_{t+K}]$, in parentheses where reported). *Eff. rank* is the effective rank measure introduced in (Roy & Vetterli, 2007).

Table 9: **Effective Rank and Linear Probing** of object, proprio, and action information across `PushT`, `Cube`, `Scene`, `Puzzle`. Action-prediction improves LeWAM’s proprioceptive and action awareness relative to LeWM, while SIGReg helps LeWAM preserve more object and general state information compared to the same model trained without it.

Environment	Model	SIGReg	Eff. rank	Object $\uparrow$	Proprio $\uparrow$	Action (IDM) $\uparrow$
PushT	LeWAM	Yes	42.90	0.90	0.86	0.40 (0.75)
		No	8.20	0.69	0.86	0.38 (0.73)
	LeWAM (MSE)	No	7.08	0.20	0.84	0.51 (0.74)
	LeWM	Yes	94.20	0.95	0.80	0.29 (0.71)
		No	16.23	0.00	0.00	0.00 (0.00)
Cube	LeWAM	Yes	46.60	0.99	0.67	0.37 (0.66)
		No	8.00	0.99	0.70	0.37 (0.68)
	LeWM	Yes	108.10	0.99	0.58	0.17 (0.55)
Scene	LeWAM	Yes	48.90	0.48	0.78	0.39 (0.74)
		No	12.50	0.40	0.77	0.43 (0.74)
	LeWM	Yes	111.10	0.40	0.37	0.07 (0.35)
Puzzle	LeWAM	Yes	46.50	0.37	0.70	0.33 (0.69)
		No	5.80	0.22	0.66	0.31 (0.68)
	LeWM	Yes	97.40	0.60	0.42	0.16 (0.48)

C.4 VISUAL RECONSTRUCTIONS AND OPEN-LOOP ROLLOUTS


Figures 13 and 14 align recorded frames with pixel-decoded representations and latent predictions on `Push-T`, `Cube`, `Scene`, and `Puzzle`. Framewise reconstructions encode and decode each recorded observation independently. Open-loop predictions start from the initial observation context and advance using the same action sequence previously executed by LeWAM. These recorded actions are fixed inputs to each model’s dynamics branch; LeWAM’s action flow branch does not generate new actions during the imagined rollout. Neither model receives intermediate recorded observations. Each prediction covers a five-step block. We sample the 51-frame clips every five frames, displaying indices relative to the clip start. The selected episode is the lowest-ID episode among the three supplied examples for each task.

The pixel decoders are trained post hoc on frozen world models using 150,000 sampled frames for 20 epochs; the displayed episodes are excluded from decoder training. These diagnostic clips use 50-epoch LeWAM checkpoints on `Cube`, `Scene`, and `Puzzle`, and an epoch-16 checkpoint on the full `Push-T` dataset. LeWM uses our 15-epoch checkpoints from its official implementation on the three 3D tasks and its released weights on `Push-T`. These checkpoint choices apply to the visual diagnostic.

On `Push-T`, both models produce visually plausible pusher motion and block translation and rotation under the supplied actions (Figure 13), providing qualitative evidence of useful action-conditioned dynamics in both models. On the three 3D manipulation tasks, `Cube`, `Scene`, and `Puzzle`, LeWAM preserves clearer details of the robot arm, gripper, and manipulated objects in its reconstructions and decoded predictions (Figures 13 and 14). The differences are especially visible in task-relevant regions. This is consistent with joint action flow-matching supervision encouraging control-relevant representations and dynamics, even though no predicted actions are used during these rollouts. The reported model sizes are comparable: approximately 16.9M parameters for LeWAM and 15M for LeWM (Maes et al., 2026b). The contribution of action supervision is not isolated by this comparison, since the encoders, training schedules, and fitted pixel decoders also differ.

(a) Push-T

Episode 1235

**(b) Cube**

Episode 638




Figure 13: **Aligned Visual Sequences on Push-T and Cube.** Each task shows recorded ground truth, framewise reconstructions for LeWM and LeWAM, and decoded open-loop predictions for both models. Columns use the same relative frame indices; predictions follow the recorded actions.




Figure 14: **Aligned Visual Sequences on Scene and Puzzle.** Rows show recorded ground truth, LeWM and LeWAM framewise reconstructions, and their decoded open-loop predictions. Columns use the same relative frame indices and the same recorded action sequence for both prediction rows.

C.5 MULTIMODAL ACTION DISTRIBUTIONS

To evaluate multi-route action distribution modeling, we collect expert demonstrations on a two-door navigation task with a fixed goal and perturbed initial positions. The dataset evenly assigns 50% of trajectories to the upper door and 50% to the lower door independently of starting perturbations, yielding two distinct valid routes for identical start-goal pairs. We train LeWAM using (1) our default flow-matching action head and (2) a deterministic MSE action head on the same data. During closed-loop policy rollouts (Figure 15), the flow head follows both routes, whereas the deterministic MSE head concentrates on the lower route and produces some wall collisions. This demonstrates that modeling multimodal action distributions provides comprehensive candidate coverage for planning in environments with topological or multi-route ambiguities.




Figure 15: **Two-Door Demonstrations and Action-Head Rollouts.** Expert demonstrations (a), executed trajectories from the flow-matching action head (b), and executed trajectories from the deterministic MSE action head (c). Black segments mark walls, and the green cross marks the goal. Blue and orange indicate upper- and lower-door trajectories; dashed red paths end in collision.

C.6 SIMULATION TASK-COMPLETION RESULTS

Table 10 reports the per-task results and deviations behind Figure 8(d).

Table 10: Task completion. Mean success rates (%) $\pm$ standard deviation across three evaluation seeds, with 50 episodes per seed. Bold marks the better result.

Method	DrawerCleanup	Transport	ToolHang	Average
Diffusion Policy (Chi et al., 2023)	54.7 $\pm$ 1.2	45.3 $\pm$ 9.0	84.7 $\pm$ 3.1	61.6
LeWAM	52.7 $\pm$ 5.0	61.3 $\pm$ 9.9	80.0 $\pm$ 3.5	64.7

D PHYSICAL-ROBOT SETUP AND EVALUATION

This section documents the physical-robot experiments in Section 4.4.

D.1 PLATFORM AND DATA COLLECTION

Demonstrations are collected by joint-space teleoperation from a reBot Arm 102 leader to a reBot B601-DM follower through LeRobot (Cadene et al., 2026). The follower uses impedance control with gravity feedforward. One proficient operator collects successful demonstrations without intentional exploration. Between episodes, both arms return to the calibration zero pose before recording resumes.

A ZED-Mini records the scene at 672×376 pixels per eye, and a wrist-mounted Logitech C270 records at 640×360 pixels. The model uses the left scene view and wrist view, both cropped to 256×256 pixels; the right stereo view is recorded but unused. At 30 Hz, a Raspberry Pi 5 stores the three video streams, seven-dimensional joint targets as actions, and follower position, torque, and velocity as a 21-dimensional state.

Table 11: Datasets for the physical-robot tasks in Section 4.4.

Task	Episodes	Frames	Duration	Task Description
Pick and Place T	200	52,297	29 min	"Pick up the T block and place it onto the outline."
Push-T	100	61,996	34 min	"Push the T block onto the outline."
Tape Stack	100	23,726	13 min	"Stack the green roll of tape on the white roll of tape."

D.2 MODEL IMPLEMENTATION

The LeWAM models in Section 4.4 use the state MSE and action flow matching objectives from Section 3.1. In the frozen V-JEPA 2.1 variant, attentive pooling and a BatchNorm MLP reduce each camera frame to one summary token. Proprioceptive state and previous actions are linearly projected from raw vectors to the model dimension. Learned queries initialize the future-camera and action streams.

The real world LeWAM variants take two video context frames and predict three future frames, all sampled at 2.5 Hz. They use model width 512 and depth eight, and train with AdamW at batch size 256, using gradient accumulation when needed. Both variants use a ViT-B visual backbone. The end-to-end model trains its visual encoder from scratch; the frozen variant initializes its encoder from V-JEPA 2.1 (Mur-Labadia et al., 2026) and keeps it fixed. For exact reproduction, check the published configuration on the project page.

For Table 2, batch-1 GFLOPs include all inference steps for one full action chunk: 100 DDPM steps for Diffusion Policy (Chi et al., 2023), ten Euler steps for $\pi_{0.5}$ (Physical Intelligence et al., 2025), and ten Euler steps for each physical-robot LeWAM variant.

D.3 PHYSICAL-ROBOT BASELINE CONFIGURATIONS

For the comparison in Section 4.4, both baselines use the LeRobot implementations (Cadene et al., 2026), adapted to the scene and wrist cameras and seven-dimensional joint-target actions. Each policy returns 36 future actions and executes the first 18 at 30 Hz before requesting a new chunk, matching the LeWAM variants. The returned and executed chunks span 1.2 and 0.6 seconds, respectively. The demonstrations and evaluation procedure are described in Appendices D.1 and D.4, with 30 trials per task and method.

The `Pick and Place T` and `Push-T` training configurations resize both camera streams to 256×256 and retain seven joint torques from the recorded proprioceptive state. Actions and torques are normalized by mapping their training-set first and 99th percentiles to -1 and 1 , respectively. Training examples start every three recorded frames; the observation offsets and action targets retain their 30 Hz timing. We hold out 10% of episodes using split seed zero. The configurations specify 20,000 optimizer updates with effective batch size 256, a 500-update linear learning-rate warmup, and cosine decay. Validation runs every 200 updates; the trainer saves both the final checkpoint and the checkpoint with the lowest validation loss.

- **Diffusion Policy** (Chi et al., 2023). The robot implementation uses a separate ImageNet-initialized ResNet-18 with spatial softmax for each camera and a conditional one-dimensional U-Net with channel widths (336, 672, 1344). It conditions on two consecutive image and torque observations. Images are randomly cropped to 243×243 during training and center-cropped at evaluation. The internal diffusion horizon is 40 action positions, aligned from one step before the current observation; the policy returns positions corresponding to the next 36 actions. This 40-step internal sequence satisfies the U-Net’s temporal downsampling constraint. Training uses noise-prediction MSE with a squared-cosine schedule, and inference uses 100 DDPM steps. AdamW uses learning rate 10^{-4} , betas (0.95, 0.999), and weight decay 10^{-6} . Exponential moving averaging with decay 0.999 supplies the weights used for validation and saved checkpoints.
- $\pi_{0.5}$ (Physical Intelligence et al., 2025). The policy conditions on the current scene and wrist images together with the task instruction and current torque vector. Its image frontend resizes the inputs to 224×224 . The wrapper discretizes the normalized torque values into 256 bins and includes

them in the language prompt. The action expert predicts a 36-action chunk by flow matching, using ten Euler steps at inference. The expert-only configuration freezes the vision and language backbone while training the action expert and action/time projections. Its task configurations use AdamW with learning rate 2.5×10^{-5} and weight decay 10^{-4} , and gradient accumulation to reach effective batch size 256.

D.4 CLOSED-LOOP EVALUATION

For Section 4.4, the robot sends its observation window to a remote inference machine, receives an action chunk, and executes it at 30 Hz under the receding-horizon schedule. A human evaluator labels each episode as success or failure, using the same criteria used during teleoperation. Each method is evaluated in 30 trials per task.

D.5 REPRESENTATION DIAGNOSTICS AND EVALUATION

For visualization, we fit a pixel decoder d_ψ to frozen encoder outputs $\mathbf{z}_t = \text{Enc}(\mathbf{o}_t)$. Learnable per-patch queries use cross-attention to map these representations to pixel patches. Only the decoder parameters ψ are trained, using per-pixel MSE over training frames:

$$\mathcal{L}_{\text{dec}}(\psi) = \mathbb{E}_{\mathbf{o}_t} [\text{MSE}(d_\psi(\text{Enc}(\mathbf{o}_t)), \mathbf{o}_t)] . \quad (12)$$

The encoder and latent dynamics remain frozen. Framewise reconstructions are $d_\psi(\mathbf{z}_t)$, and decoded rollout predictions are $d_\psi(\hat{\mathbf{z}}_t)$.

Figure 16 compares both encoder variants on the same held-out Pick and Place T episode from the scene camera. Each strip shows recorded video, framewise reconstructions, and decoded autoregressive predictions. Latent rollouts start from the initial observation context and advance under recorded actions, without intermediate ground-truth frames or newly generated actions. We display ten frames at indices 0, 3, 6, 9, 12, 15, 18, 21, 24, and 26, spanning 10.4 seconds at the 2.5 Hz model sampling rate.




Figure 16: Matched visual reconstructions and recorded-action latent rollouts for the scene camera. Within each model, rows show the ground-truth (GT) video, independent frame reconstructions, and decoded autoregressive predictions.

E DISCUSSION OF THE ROLE OF JOINT PREDICTION

This section discusses the role of joint prediction in representation learning, action-conditioned dynamics, and planning, complementing Sections 4 and 3.2. The schematics and idealized conditional-risk analysis motivate possible mechanisms, whose causal contributions have not been isolated by our experiments.

Table 12 complements Section 2 with a comparison of training requirements and prediction capabilities.

Table 12: **Modeling and Control Capabilities.** FDM denotes forward dynamics conditioned on clean actions; Policy denotes learned action prediction. Entries describe reported training recipes or the goal-reaching configurations used here.

Method	FDM	Policy	Joint FDM and Policy	End-to-end from scratch
<i>World models</i>				
DINO-WM (Zhou et al., 2024)	✓	✗	✗	✗
V-JEPA 2-AC (Assran et al., 2025)	✓	✗	✗	✗
LeWM (CEM) (Maes et al., 2026b)	✓	✗	✗	✓
<i>World action models</i>				
UVA (Li et al., 2025)	✓	✓	✗	✗
UWM (Zhu et al., 2025)	✓	✓	✗	✗
DreamZero (Ye et al., 2026)	✗	✓	✗	✗
Fast-WAM (Yuan et al., 2026)	✗	✓	✗	✗
Cosmos 3 (NVIDIA, 2026)	✓	✓	✗	✗
LeWAM (ours)	✓	✓	✓	✓
<i>Goal-conditioned policy</i>				
LeWM (IDM) (Nguyen et al., 2026)	✓	✓	✗	✗
GCBC (Ghosh et al., 2021)	✗	✓	✗	✗

Joint requires clean-action FDM and policy losses on the same sample in one training pass. From scratch includes visual encoding and prediction without pretrained or frozen components. UWM uses a frozen pretrained variational autoencoder even when its robot backbone is trained from scratch.

E.1 REPRESENTATION AND DYNAMICS

Representation Shaping. SIGReg (Balestriero & LeCun, 2025; Maes et al., 2026b) regularizes the latent marginal; action supervision encourages control-relevant distinctions (Figure 17). If visual history changes the conditional action velocity, discarding it increases optimal action regression risk. Latent prediction adds pressure to retain action-conditioned consequences. The diagram is conceptual: action prediction alone does not rule out every form of collapse.




Figure 17: Conceptual roles of SIGReg and action flow matching in shaping the latent space.

Conditional Action Risk. In Equation 2, squared flow matching makes the optimal velocity the conditional mean of $\mathbf{a}_j - \boldsymbol{\epsilon}_j$. Let $\mathcal{I}_j$ include the noisy action prefix, flow time, and other conditioning, but exclude visual history $\mathbf{z}_{\leq t}$. With $\mathcal{R}_{\text{FM}}$ denoting optimal velocity MSE, the conditional projection identity gives

$$\begin{aligned} & \mathcal{R}_{\text{FM}}(\text{constant}) - \mathcal{R}_{\text{FM}}(\mathbf{z}_{\leq t}) \\ &= \frac{1}{C} \sum_{j=t}^{t+C-1} \mathbb{E}[\text{MSE}(\mathbb{E}[\mathbf{a}_j - \boldsymbol{\epsilon}_j \mid \mathcal{I}_j, \mathbf{z}_{\leq t}], \mathbb{E}[\mathbf{a}_j - \boldsymbol{\epsilon}_j \mid \mathcal{I}_j])] \geq 0. \end{aligned} \quad (13)$$

When visual history changes the conditional velocity, a constant code incurs higher optimal risk. This pressure can vanish if other inputs already determine the velocity. Thus action prediction cannot exclude partial collapse; LeWAM also uses SIGReg to regularize the global latent distribution.

E.2 PLANNING WITH LEARNED ACTION PROPOSALS




Figure 18: **Action Proposals and Predicted Consequences.** (a) Smooth rollouts; (b) candidate rollouts branch at contact; (c) LeWAM pairs learned action proposals with predicted outcomes. Dashed links denote search; the solid link denotes joint prediction. This schematic illustrates finite-budget search, not a universal CEM failure.

Behavior and Consequence. The LeWM (Maes et al., 2026b) baseline learns action-conditioned latent dynamics without an action-prediction objective and uses CEM (Rubinstein, 1999) for control. Joint prediction additionally supplies a learned action distribution for planning proposals. A contact event can switch the active mode, so nearby actions yield distinct outcomes (Pang et al., 2023). If sampled plans miss useful contact, an object-based goal cost provides little ranking signal (Yamada et al., 2024). CEM is derivative-free and can cross nonsmooth transitions, but a finite population cannot favor a useful branch it has not sampled; more iterations increase rollout cost. Learned proposals can improve contact-rich planning under constrained budgets (Brudermüller et al., 2025). Figure 18 illustrates LeWAM’s use of learned proposals and predicted consequences. The planning objective is defined in Section 3.2, and the refinement procedure is specified in Appendix A.1.

F AUTHOR CONTRIBUTION STATEMENT

All coauthors contributed to discussions, provided input on all aspects of the project, and helped with writing and editing.

Minghao initiated the project and led its technical development. He proposed the general architecture with action flow matching and MoT and the initial version of gradient-based planning. He also constructed the benchmarks and wrote the main part of the paper.

Tavis designed the MoT attention mechanism, optimized the planning algorithm to use batched gradients, and refined the methodology through extensive experimentation. He also optimized the code for release. Tavis and Minghao completed the simulation experiments reported in the main paper.

Eryk designed and carried out the entire physical robot study. He built the hardware and teleoperation pipeline, adapted the model for real-robot deployment, collected all demonstrations, trained all real-robot models, and ran the physical evaluations. He also wrote the real-robot experiments section and produced its figures.

Randall initiated the LeWAM project and provided overall project leadership, supervision, and management. He guided the project’s branding and contributed extensively to manuscript writing and editing. He also confirmed the paper’s final title.